

Advanced DevOps / SRE / AWS Questions

Q1. How do you design high-availability architecture across multiple regions?

Answer:

- Multi-region EKS clusters
- Route53 latency-based routing
- DB replication (Aurora Global / DynamoDB Global Tables)
- S3 cross-region replication
- Multi-AZ subnets
- Auto-scaling groups
- Health checks + Failover policies

Q2. Explain how you implemented Disaster Recovery.

Answer:

- Automated backups (RDS/Aurora snapshots)
- DynamoDB cross-region replication
- EKS cluster backups
- Replicated S3 buckets
- Failover routing with Route53
- Weekly DR drills to validate RTO/RPO

Q3. How do you evaluate cost optimization?

Answer:

- Identify idle instances
- Use instance right-sizing
- Switch to Graviton instances
- Use Spot + On-Demand mix
- Clean unused EBS, snapshots, ELBs
- Lifecycle policies for S3
- Enable Compute Savings Plans

Q4. How do you ensure zero downtime deployments?

Answer:

- Canary / Blue-Green
- Readiness + liveness probes
- Auto-rollback in Harness
- Pod disruption budgets
- Rolling updates with controlled maxSurge and maxUnavailable

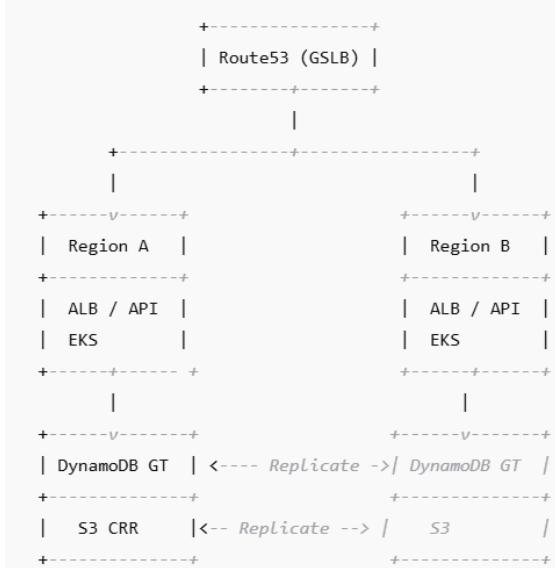
Q5. How do you secure Kubernetes?

Answer:

- RBAC control
- Pod Security Policies / OPA Gatekeeper
- Image scanning (JFrog/Qualys)
- Secrets encryption using AWS Secrets Manager
- Network policies
- Disable privileged containers

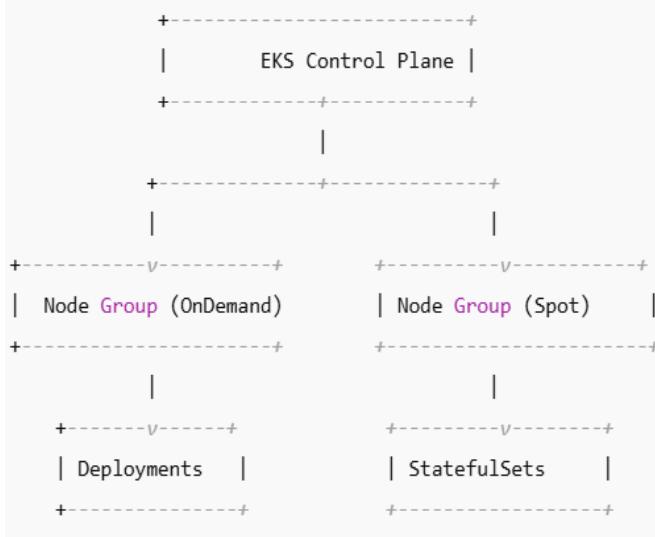
1. Design a Multi-Region, Active-Active Architecture

lua

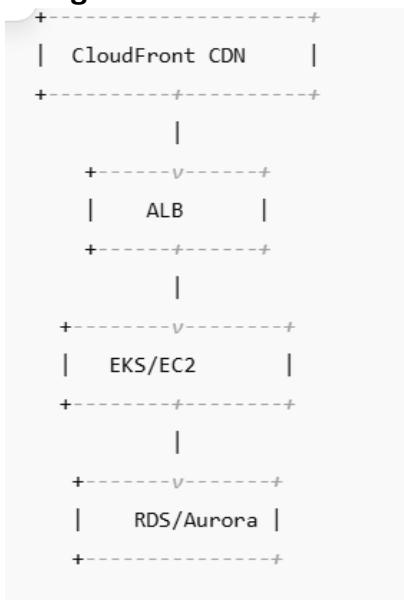


2. Design a Highly Available EKS Cluster

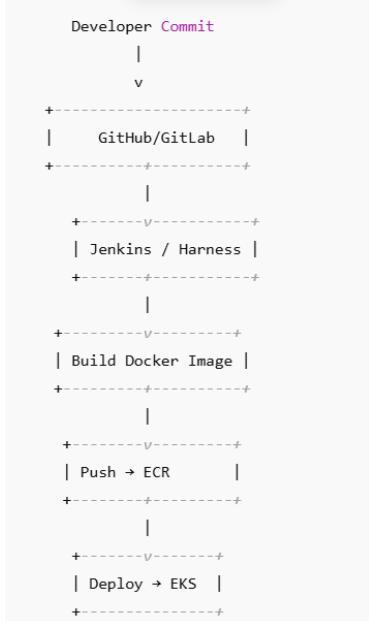
sql



3. Design a Global CDN-Based Architecture



4. DevOps CI/CD Architecture



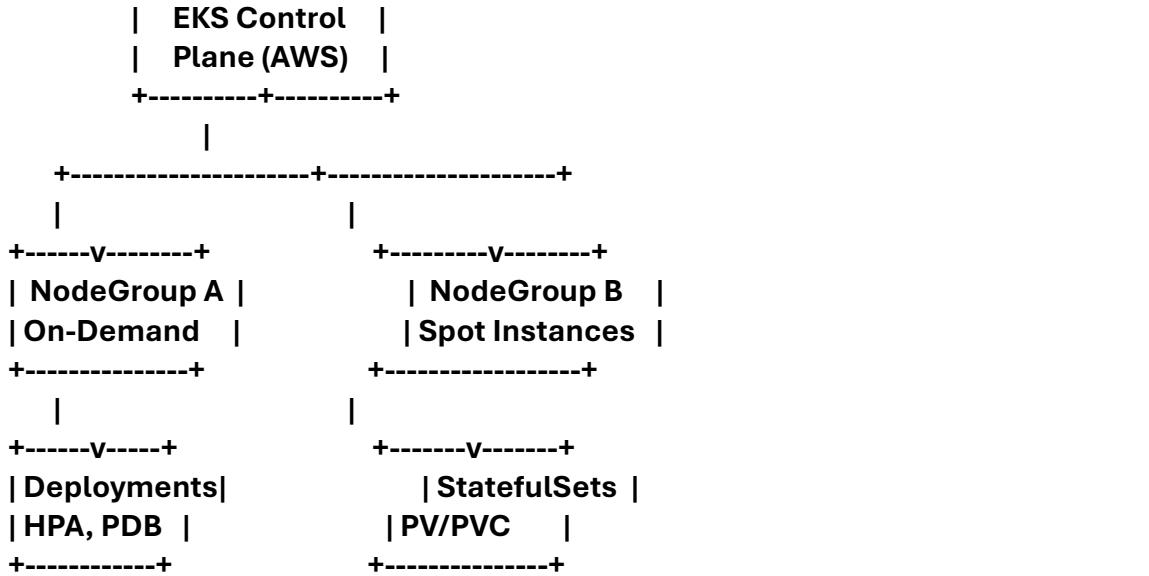
ARCHITECTURE 2 — Highly Available EKS Cluster (Production-Ready)

Goal:

Build a resilient EKS cluster that supports zero downtime, automatic scaling, and safe deployments.

Diagram





Step-by-Step Design

Step 1: Control Plane

Managed by AWS:

- HA across 3 AZs
 - Auto-scaled by AWS
 - Use version lifecycle: EKS 1.27+ recommended
 - Logging enabled: audit, api, scheduler, authenticator
-

Step 2: Node Groups

Create 3 node groups:

1. On-demand nodes:
 - Core services
 - System pods
 - Ingress controllers
 2. Spot nodes:
 - Non-critical workloads
 - Autoscaled
 3. GPU node group (optional)
 - AI/ML workloads
-

Step 3: Networking

AWS VPC CNI + Security best practices:

- Assign IAM roles to service accounts
- Pod-level security via:
 - Network Policies
 - Istio policies
 - Calico (optional)

Step 4: Autoscaling

1. **HPA: CPU/Mem/custom metrics**
 2. **Cluster Autoscaler:**
 - o Adds nodes when pods are Pending
 3. **Karpenter (new option):**
 - o More powerful
 - o Spot instance orchestration
-

Step 5: Storage

- **gp3 EBS volumes for PVCs**
 - **EFS for shared content**
 - **fsx for high-performance workloads**
-

Step 6: Deployment Safety

Use:

- **Readiness & liveness probes**
 - **PodDisruptionBudgets (PDB)**
 - **RollingUpdate strategy**
 - **Blue-Green via Argo Rollouts / Harness**
-

Step 7: Security

- **RBAC**
 - **OPA Gatekeeper**
 - **AWS Secrets Manager**
 - **KMS encryption**
 - **mTLS via Istio**
-

Step 8: Monitoring

- **Prometheus + Grafana**
 - **CloudWatch Container Insights**
 - **Jaeger / X-Ray for tracing**
-

ARCHITECTURE 6 — Multi-Account AWS Landing Zone (Enterprise-Grade)

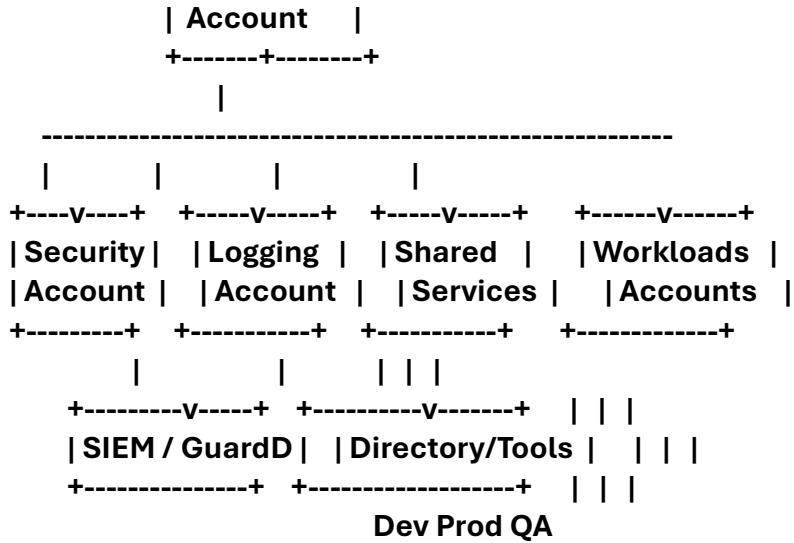
Goal:

A secure enterprise AWS foundation using multi-account isolation, SCP controls, centralized logging, shared services, and secure networking.

High-Level Diagram

AWS ORGANIZATION





🛠 Step-by-Step Design

Step 1: Set up AWS Organizations

- Create root organization
- Define OU (Organizational Units):
 - *Security OU*
 - *Infrastructure OU*
 - *Production OU*
 - *Sandbox OU*

Why?

Logical grouping for applying Service Control Policies (SCPs).

Step 2: Mandatory Accounts

1. Management Account

- Only for billing and org management
- No workloads allowed

2. Security Account

- GuardDuty
- Security Hub
- IAM Access Analyzer
- Inspector
- Centralized IAM Role creation

3. Logging Account

- Centralized CloudTrail
- VPC Flow Logs
- Lambda function logs
- S3 with retention

4. Shared Services Account

- Directory Services

- DNS
- Jenkins/Harness
- Jump/Bastion Hosts
- Transit Gateway

5. Workload Accounts

- Dev
- QA
- Prod
- ML
- Analytics

Why?

Blast-radius reduction and security isolation.

Step 3: Service Control Policies (SCPs)

Examples:

- Deny disabling CloudTrail
 - Deny public S3 buckets
 - Deny IAM policy wildcards
 - Restrict which regions can be used
-

Step 4: Centralized Logging Architecture

Account A --> CloudTrail -->|

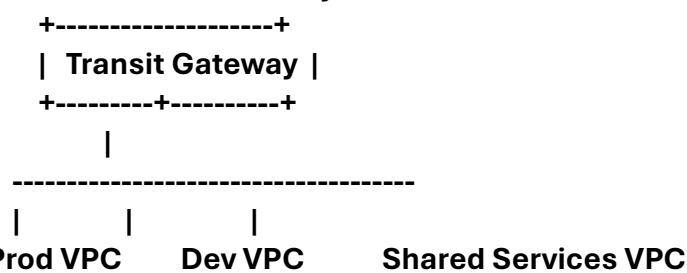
Account B --> Flow Logs ---->|--> Logging Account (S3 + OpenSearch)

Account C --> App Logs -----|

All logs are immutable, versioned, and encrypted with KMS.

Step 5: Networking Setup

Use AWS Transit Gateway to connect all VPCs:



Benefits:

- Dedicated VPC per account
 - Centralized inspection (firewalls)
 - No overlapping CIDRs
-

Step 6: Identity & Access

- IAM Identity Center / SSO
- RBAC per account
- MFA enforced

- **Cross-account IAM roles**
-

Step 7: Security Foundation

- **GuardDuty enabled org-wide**
 - **Security Hub with CIS Level 1 & 2**
 - **AWS Config rules enabled**
 - **S3 access logs mandatory**
 - **KMS for encryption everywhere**
-

Step 8: Cost Optimization

- **Consolidated billing**
 - **Reserved Instances for shared services**
 - **Budgets and anomaly detection**
-

Use Cases

- ✓ **Fortune-500 companies**
- ✓ **Regulated industries**
- ✓ **Multi-team large orgs**

ARCHITECTURE 8 — Serverless Event-Driven Platform (Kinesis / Lambda)

Goal:

Process millions of events/minute with elastic scaling and minimal ops overhead.

Diagram

Clients → API Gateway → Kinesis Stream → Lambda Consumers → DynamoDB/S3
|
Firehose → S3 / Redshift

Steps

Step 1: API Ingestion Layer

Use API Gateway:

- Rate limiting
 - AuthZ/A (JWT, IAM)
 - Request validation
-

Step 2: Kinesis Data Stream

- Shards handle scaling
- Partition key determines ordering
- 1MB/sec or 1000 records/sec per shard

Scaling:

Add shards → horizontal scaling

Resharding is online

Step 3: Lambda Consumers

Lambda pulls batches from Kinesis:

- Batch size = 100–500
 - Retries enabled
 - DLQ = SQS
 - Timeout optimized (~1–2 sec)
-

Step 4: Firehose for Heavy ETL

Use Firehose to deliver:

- Parquet data → S3
 - JSON → Redshift
 - Transformed results via Lambda
-

Step 5: Storage Layer

- DynamoDB for low-latency key-value
 - S3 for raw data lake
 - Redshift/Athena for analytics
-

Step 6: Observability

- CloudWatch Log Insights
 - X-Ray for end-to-end latency
 - Lambda concurrency alarms
-

Step 7: Cost Optimization

- Use on-demand Kinesis scaling
 - Efficient partition keys
 - Compress S3 objects
-

ARCHITECTURE 12 — IoT Platform Architecture (MQTT + IoT Core + Kinesis + DynamoDB)

Goal:

Support millions of IoT devices sending telemetry in real-time with secure ingestion, processing, device shadows, and two-way communication.

Diagram

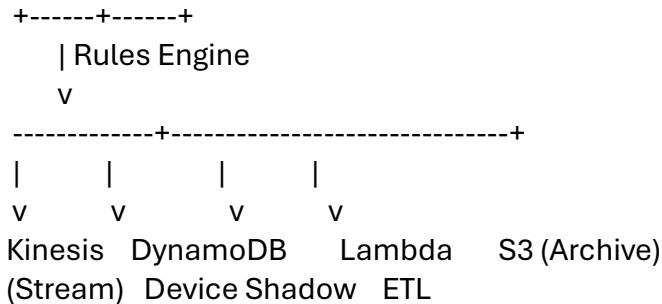
IoT Devices (MQTT)

|

v

+-----+

| AWS IoT Core|



🛠 Step-by-Step Design

Step 1: Device Connectivity

Devices connect via **MQTT over TLS 1.2**

Auth:

- X.509 certificates
- Cognito tokens (optional)

All traffic encrypted.

Step 2: Ingestion (IoT Core Broker)

IoT Core broker receives MQTT messages and triggers **Rules Engine**.

Step 3: Rules Engine Routing

Routes incoming messages to:

- DynamoDB (for Device Shadow)
- Kinesis (for streaming analytics)
- Lambda (for processing commands)
- S3 (for cold storage)

Step 4: Device Shadow System

Manages current state of device:

- Desired state
- Reported state
- Delta state

Used for remote control:

- Turn fan ON/OFF
- Update thermostat
- Push firmware

Step 5: Real-time Processing

Kinesis → Lambda

- Aggregates temperature/humidity
- Detects anomalies
- Pushes alerts

DynamoDB Device Shadow Updates

- Low-latency device state
 - Fast read/write
-

Step 6: Batch Analytics

S3 data processed through:

- Athena
 - EMR
 - Glue
 - Redshift
-

Step 7: Device Commands

Cloud → IoT Core → Device

Via MQTT with QoS.

Step 8: Security

- Mutual TLS
- IAM policies per certificate
- Device revocation list
- JITR/JITP certificate provisioning

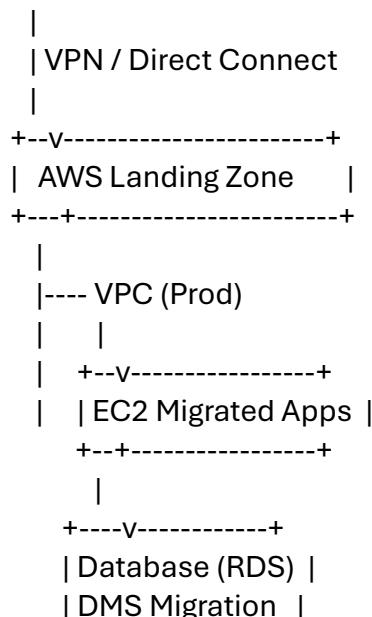
ARCHITECTURE 14 — Lift & Shift Migration Architecture (On-Prem → AWS)

Goal:

Migrate legacy workloads with minimal redesign using VPN/Direct Connect, Database Migration Service (DMS), and VM imports.

Diagram

On-Prem DC



+-----+

☒ Step-by-Step Migration Strategy

Step 1: Network Connectivity

Options:

- Site-to-Site VPN (quick)
- Direct Connect (stable, low latency)

Setup:

- Transit Gateway to attach multiple VPCs
- Route filtering
- BGP for dynamic routing (DX)

Step 2: Assessment

Tools:

- AWS Application Discovery Service
- Migration Evaluator

Identify:

- CPU usage
- Storage needs
- Licensing
- Dependencies

Step 3: Server Migration

Use **AWS Application Migration Service (MGN)**

Steps:

1. Install agent on on-prem servers
2. Continuous replication to AWS
3. Launch test instances
4. Cutover window selected
5. Switch DNS

Step 4: Database Migration

Use **AWS DMS**:

- Minimal downtime
- Continuous replication
- Schema conversion via SCT

Step 5: Storage Migration

Use:

- AWS Snowball for large TB/PB
- AWS DataSync for continuous sync

Step 6: User Identity Integration

- Connect AD → AWS AD Connector
 - Single Sign-On
-

Step 7: DR Setup

- Backups
 - Cross-region replication
 - RPO/RTO definitions
-

Step 8: Optimization

Post-migration:

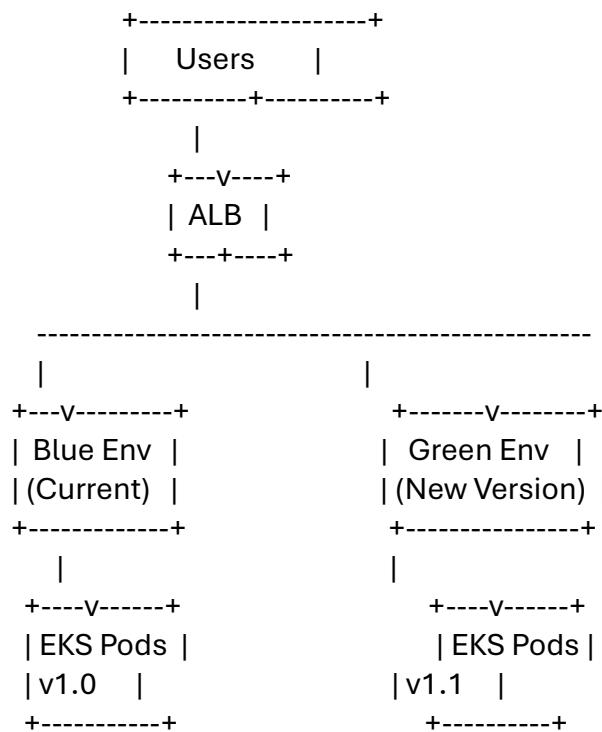
- Convert EC2 → ECS/EKS
 - Move DB to Aurora
 - Add auto-scaling
 - Add caching layer
-

ARCHITECTURE 16 — Zero-Downtime Blue/Green Deployment Architecture (EKS + ALB + CI/CD)

Goal:

Release applications with **zero downtime**, instant rollbacks, and low-risk deployments.

✿ High-Level Diagram



🛠 Step-by-Step Design

Step 1: Two Identical Environments

- **Blue environment** = running production version
- **Green environment** = new version being deployed

Both deployed in:

- Separate EKS namespaces
- Or separate EKS deployments
- Or separate clusters (large enterprises)

Step 2: ALB Target Groups

ALB configured with:

- **TG-Blue** → Blue pods
- **TG-Green** → Green pods

Routing = 100% to Blue initially.

Step 3: CI/CD Pipeline

Pipeline steps:

1. Build Docker image
2. Security scan (SNYK/Trivy)
3. Deploy to Green environment
4. Run smoke tests
5. Validate logs, error rates, latencies

Step 4: Traffic Switch

Once green environment validated:

- Shift ALB routing from Blue → Green
- Begins with partial:
 - 10% traffic
 - 25%
 - 50%
 - 100%

Or immediate full cutover.

Step 5: Instant Rollback

If errors or latency spike:

- Route back to **Blue**
- Green scaled down but not removed
- Debug safely

Rollback time: < 5 seconds

Step 6: Cleanup

Once stable:

- Destroy Blue
- Mark Green as new Blue
- Next deployment → new Green environment

Step 7: Monitoring & Observability

- Compare both versions side-by-side
- Key metrics:
 - 5xx errors
 - p95 latency
 - CPU/memory for pods
 - Logging anomalies

Step 8: Best Practices

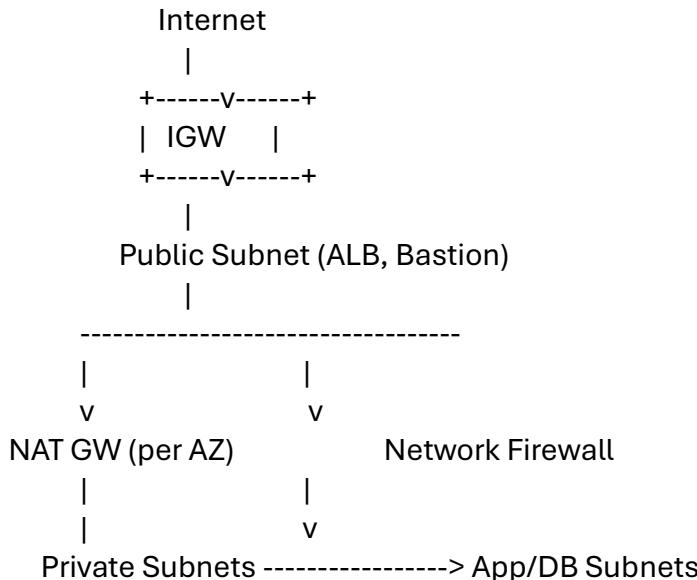
- Don't share DB write schemas
- Use read replicas for migration testing
- Keep both environments identical
- Test failure injection on green

ARCHITECTURE 18 — VPC with Public/Private Subnets + NAT + Bastion + Firewalls

Goal:

A secure AWS network where private servers have no public exposure, outbound internet is controlled, and access is tightly limited.

Diagram



🛠 Step-by-Step Design

Step 1: VPC & Subnet Layout

Create:

- 3 AZs
- 3 public subnets

- 3 private subnets
- 3 DB subnets

CIDR example:

- VPC: 10.0.0.0/16
- public: 10.0.1.0/24
- private: 10.0.2.0/24
- db: 10.0.3.0/24

Step 2: Internet Gateway + Public Subnets

Public subnets host:

- Bastion host
- ALB
- NAT Gateway

Routes:

0.0.0.0/0 → IGW

Step 3: NAT Gateway in Each AZ

Private subnets route outbound via NAT:

0.0.0.0/0 → NAT-A

NAT is **highly available** when placed per-AZ.

Step 4: Private Subnets

Private workloads:

- EKS worker nodes
- ECS tasks
- EC2
- Lambda in VPC

No public internet exposure.

Step 5: Database Subnets

For RDS/Elasticache:

- No outbound
- No inbound except from app subnets
- Encrypted in transit

Step 6: Bastion Host

Access via:

- SSH session manager (preferred)
- MFA authentication
- IAM role

Security:

- Restricted IP
- Logging via Systems Manager

Step 7: Network Firewall

Use **AWS Network Firewall**:

- Egress filtering
- Domain filtering
- Block malicious outbound traffic

Step 8: VPC Endpoints

For S3, DynamoDB, ECR:

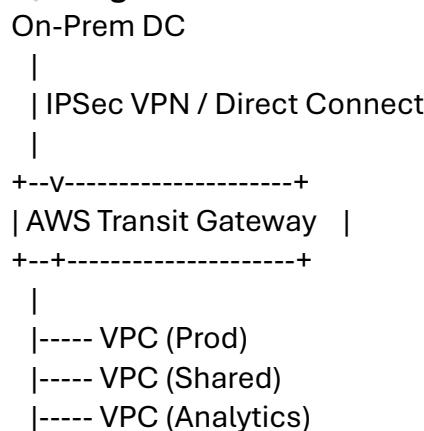
- Reduce NAT cost
- Increase security
- No public internet required

ARCHITECTURE 19 — Hybrid Cloud Architecture (On-Prem ↔ AWS)

Goal:

Connect on-prem data centers with AWS securely to form a hybrid environment for workload extension.

Diagram



Step-by-Step Design

Step 1: Network Connectivity

Choose:

- **VPN** (cheap, fast to configure)
- **Direct Connect** (stable, dedicated)

Best: Use **DX with VPN backup**.

Step 2: Transit Gateway

Connects:

- On-prem routers
- Multiple VPCs
- Virtual appliances

- Shared services

Route domain isolation:

- Prod traffic
 - Dev traffic
 - External traffic
-

Step 3: Identity Integration

Options:

- AWS Managed AD
 - AD Connector
 - SAML Federation for SSO
 - IAM Identity Center
-

Step 4: Data Sync

Use:

- AWS DataSync for file servers
 - AWS Transfer for SFTP
 - RDS Migration Service
 - Database replication from on-prem
-

Step 5: Application Integration

Hybrid workloads:

- On-prem DB + AWS app tier
 - On-prem MQ + AWS microservices
 - AWS analytics processed from on-prem sources
-

Step 6: Security

- IPSec encryption
 - Route filtering
 - Security Group + NACL boundaries
 - DDoS protection via AWS
-

Step 7: DR

On-prem → AWS DR:

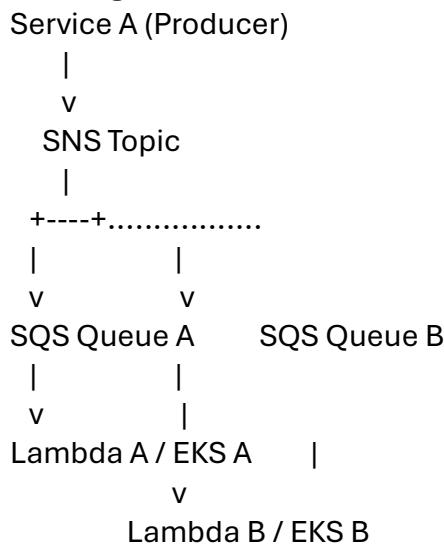
- S3 as off-site backup
 - EBS snapshots
 - RDS cross-region replication
-

ARCHITECTURE 20 — Event-Driven Microservices Architecture (SNS + SQS + Lambda/EKS)

Goal:

Loosely coupled microservices communicating via events instead of synchronous APIs.

Diagram



Step-by-Step Design

Step 1: Producer Service

Service A publishes events to **SNS Topic**.

Step 2: SNS → SQS Fan-Out

Multiple microservices can subscribe:

- SQS Queue A
- SQS Queue B
- Lambda
- HTTPS endpoints

Benefits:

- Decoupling
 - Parallel processing
 - Scalability
-

Step 3: SQS Queue Consumers

Two options:

Lambda

- Auto-scales
- Event-driven
- Low cost

EKS/ECS Services

- Always running
 - Suitable for heavy tasks
-

Step 4: Retry & DLQ

SQS → DLQ policy:

- 3 retry attempts
 - After that → DLQ
 - Alerts triggered
-

Step 5: Ordering

SQS FIFO Queue ensures:

- Exactly-once processing
 - Message order maintained
-

Step 6: Idempotency

Consumer services must:

- Generate unique message IDs
 - Avoid reprocessing repeated events
-

Step 7: Monitoring

Tools:

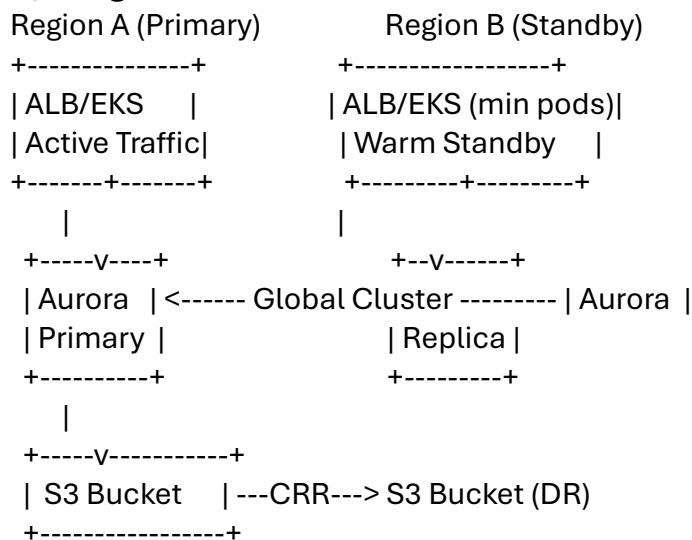
- CloudWatch Metrics
- SQS queue depth alarms
- Lambda concurrency alarms
- SNS failure notifications

ARCHITECTURE 22 — Multi-Region Active-Passive Architecture (Warm Standby)

Goal:

Provide fast DR with optimized cost by maintaining a **warm standby** in another region.

❖ Diagram



❖ Step-by-Step Architecture Design

Step 1: Primary Region Setup

Components:

- EKS cluster (full capacity)
 - ALB for routing
 - Aurora/MySQL PostgreSQL primary cluster
 - S3 for storage
 - Redis/ElastiCache
-

Step 2: Standby Region Setup

- EKS with minimal nodes
- Aurora read replica
- S3 CRR
- Offline Redis replica (optional)

Warm but not fully active.

Step 3: Data Replication

Database:

- Aurora Global DB for sub-second replication
- Failover time: **< 1 minute**

Storage:

- S3 CRR replicates objects asynchronously
 - DynamoDB Global Tables (optional)
-

Step 4: DNS Failover

Using **Route53 failover routing**:

- Primary ALB = Active
- Standby ALB = Passive
- Health checks every 30 sec

If Region A fails:

→ Route53 shifts traffic to **Standby Region B**.

Step 5: Autoscaling for DR

On failover:

- EKS CA & HPA triggered
 - Nodegroups scale from 1 → required size
 - Redis cluster spins up via automation
-

Step 6: Recovery Runbook

1. Declare failover event
2. Route53 moves traffic
3. Promote Aurora read replica
4. Redeploy pods on standby region
5. Validate endpoints
6. Announce DR-safe state

Step 7: Cost Optimization

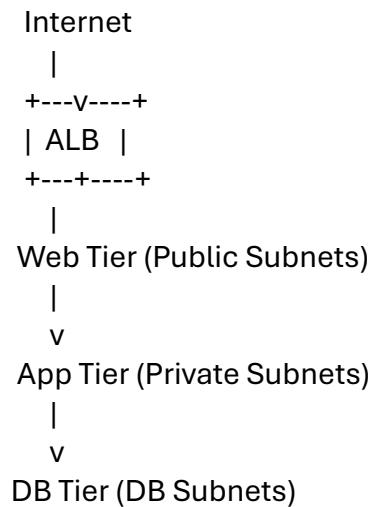
- Minimal nodes in standby
- Aurora replica costs only read capacity
- Use spot nodes in DR cluster
- S3 CRR compress objects

ARCHITECTURE 23 — Secure Multi-Tier Architecture (Web → App → DB)

Goal:

Highly compliant, layered architecture where each tier has strict access boundaries.

❖ Diagram



❖ Step-by-Step Architecture Design

Step 1: 3-Tier Network Segmentation

- **Tier 1: Web Tier**
 - ALB
 - Reverse proxies
 - API gateways
- **Tier 2: App Tier**
 - EKS apps
 - Business logic
- **Tier 3: Database Tier**
 - RDS
 - Redis
 - No public routes

Each tier in different subnets.

Step 2: Security Group Chaining

- Web SG → App SG only

- App SG → DB SG only
- DB SG → no outbound

Example rules:

ALB-SG → Allow 80/443 from public

APP-SG → Allow 8080 from ALB-SG

DB-SG → Allow 5432 from APP-SG only

Step 3: Network Access Controls

NACLs:

- Stateless
- Extra protection
- Block unknown IPs

Step 4: WAF & Shield

Deploy WAF for:

- SQL injection
- XSS
- Bot protection

Shield for DDoS.

Step 5: App Tier Architecture

EKS with:

- Pod security policies
- mTLS
- Resource quotas
- HPA

Step 6: DB Tier

- RDS Multi-AZ
- Automated backups
- Encryption at rest & transit
- Parameter group tuning

Step 7: Observability

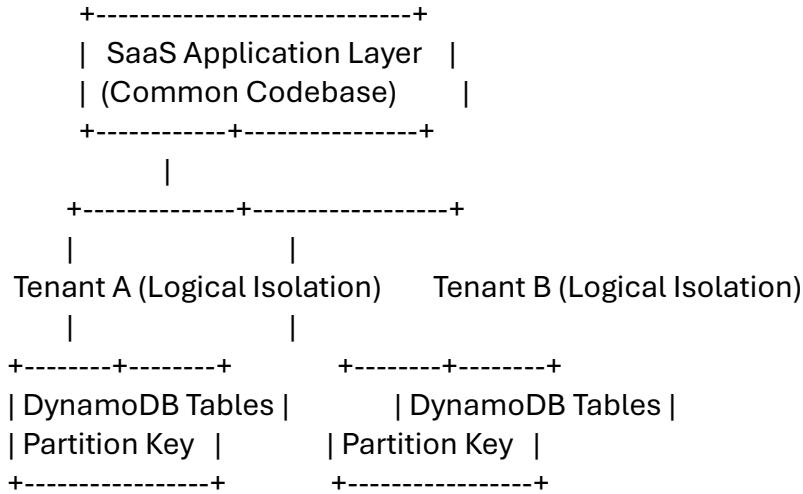
- ALB access logs
- VPC Flow logs
- X-Ray
- CloudWatch metrics

ARCHITECTURE 24 — SaaS Multi-Tenant Architecture (Tenant Isolation + Single Codebase)

Goal:

Host multiple customers (tenants) on the same application while maintaining data isolation.

❖ Diagram



❖ Step-by-Step Design

Option 1 — Pool Model (Logical Isolation, Single DB)

Single DynamoDB table with:

- Partition key = tenant_id
- IAM conditions enforce tenant-level access
- Cheaper and simpler

Option 2 — Bridge Model (Mixed)

- One shared DB
- Some resources isolated
- Ideal for mid-sized SaaS

Option 3 — Silo Model (Physical Isolation)

Separate:

- VPC
- DB
- S3 Buckets

Used for:

- High security tenants
- Financial/healthcare

Step 3: Tenant Routing

Use:

- Tenant-specific subdomain
 - customer1.myapp.com

- JWT token includes tenant_id
 - customer2.myapp.com

Step 4: Authentication

Cognito / IAM-based:

- Multi-tenant user pool
 - Tenant attributes in token

Step 5: Storage Isolation

DynamoDB isolation:

- Condition expression: `tenant_id = ${token.tenant_id}`

S3 isolation:

bucket/tenant_id/*

Step 6: Observability

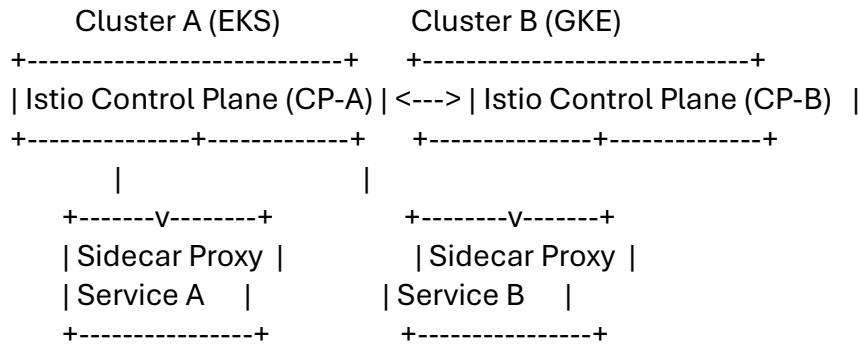
- Per-tenant dashboards
 - Per-tenant CloudWatch filters
 - Billing by tenant usage

ARCHITECTURE 30 — Service Mesh Across Multiple Clusters (Istio Multi-Cluster)

Goal:

Provide **mTLS, traffic control, cross-cluster service discovery, observability** across multiple Kubernetes clusters.

Diagram



Step-by-Step Design

Step 1: Multi-Cluster Topology Choice

Choose mode:

1. Primary-Remote

One control plane, multiple remote clusters
(Simple, centralized)

2. Primary-Primary

Each cluster has its own control plane
More complex but more resilient

Step 2: Cross-Cluster Communication

Use:

- East-West Gateway
- mTLS
- Shared root certificate authority

Service A in Cluster A can call Service B in Cluster B **securely**.

Step 3: Service Discovery

Istio replicates service registry:

- If service removed → sync to other clusters
 - Load balances requests cross-cluster
-

Step 4: Traffic Scheduling

Use DestinationRules to:

- Split traffic across clusters
- Fail over to another cluster
- Apply locality-aware routing

Example:

- 90% to local cluster
 - 10% to remote cluster
-

Step 5: Observability

Unified telemetry across clusters:

- Kiali for mesh topology
 - Prometheus for metrics
 - Jaeger for tracing
 - Grafana dashboards
-

Step 6: Zero-Trust Security

- mTLS cluster-wide
 - RBAC for services
 - JWT validation at ingress
-

Step 7: Failover

If Cluster A fails:

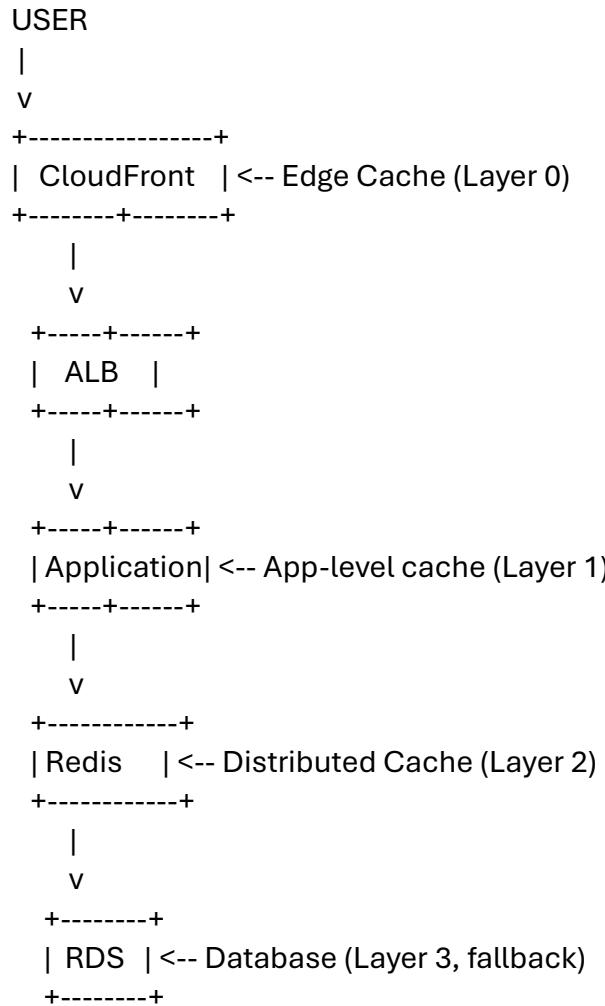
- Traffic routed to Cluster B automatically
- Control plane detects healthy endpoints
- East-West gateway handles routing

ARCHITECTURE 31 — High-Performance Caching Layer (Redis/ElastiCache + CDN + App Cache)

Goal:

Reduce database load, improve API latency, and scale read-heavy applications using multi-layer caching.

Diagram



🛠 Step-by-Step Architecture

Step 1: Layer-0 (CDN Edge Cache)

Use CloudFront for:

- Static files
- Cached API responses
- Signed URLs
- Cache invalidation

Benefits:

- ✓ Lowest latency
- ✓ Zero backend compute usage

Step 2: Layer-1 (Application Cache inside Pod/EC2)

Examples:

- In-memory dictionary
- LRU cache
- Local file cache

Benefits:

- ✓ Nanosecond access
 - ✓ No network roundtrip
 - ✗ Limited by pod lifecycle (evicted on restart)
-

Step 3: Layer-2 (Redis / ElastiCache Cluster)

Modes:

- Redis Cluster (sharded, high scale)
- Redis Primary/Replica
- Redis with Multi-AZ

Use cases:

- User sessions
 - Token verification
 - Product catalog
 - Leaderboards
 - Distributed locks
-

Key Design Considerations

✓ Eviction Policies

- LRU
- LFU
- Volatile TTL

✓ TTL Strategy

- Aggressive TTL for fast-changing data
 - Longer TTL for static data
-

Step 4: Layer-3 (RDS / Primary Store)

Fallback to DB only when Redis misses.

Step 5: Write-Through vs Write-Back vs Cache-Aside

Cache-aside (most common):

- Read → miss → fetch DB → write cache

Write-through:

- Write to cache → write to DB

Write-back:

- Write to cache only
- Async flush to DB

Step 6: Scaling Redis

- Read replicas
 - Cluster mode enabled
 - Redis I/O optimized nodes
 - Auto-failover
-

Step 7: Security

- At-rest encryption
 - In-transit encryption
 - SG restrictions
 - IAM-auth (Redis 7+)
-

Step 8: Monitoring

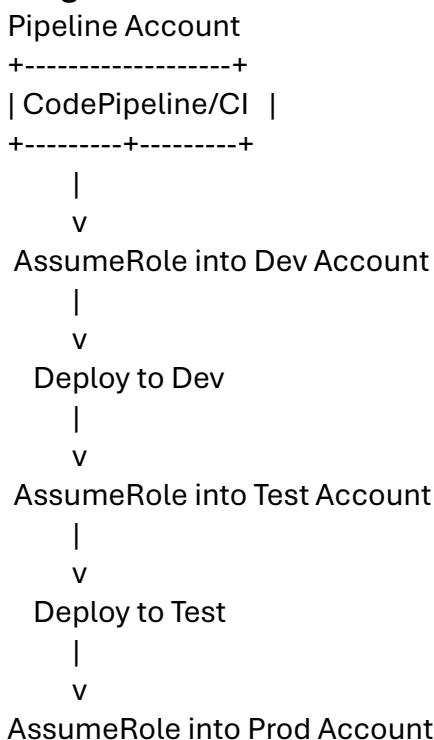
- Cache hit ratio
- Evictions
- Memory fragmentation
- Latency of GET/SET

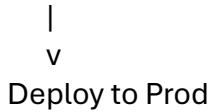
ARCHITECTURE 39 — Cross-Account CI/CD (Pipeline Account → Dev/Test/Prod)

Goal:

Centralized pipeline deploys workloads into multiple accounts securely using IAM Roles and cross-account resource access.

Diagram





🛠 Step-by-Step Architecture

Step 1: Account Structure

- CI/CD Account
- Dev Account
- Test Account
- Prod Account

Enforced by AWS Organizations.

Step 2: Cross-Account Roles

In each target account:

Role: DeployRole

Trust: Pipeline Account

Pipeline uses:

sts:AssumeRole

Step 3: CI Stage

Runs:

- CodeBuild
 - Unit tests
 - Linting
 - Image build
 - Security scanning
 - Upload artifacts to S3/ECR
-

Step 4: CD Stage

Pipeline executes:

- Deploy to dev
 - Integration tests
 - Deploy to test
 - UAT tests
 - Manual approval
 - Deploy to prod
-

Step 5: Separation of Duties

Pipeline cannot write to prod unless:

- IAM approval
 - Manual stage approval
 - Multi-level checks
-

Step 6: Auditing

CloudTrail logs capture:

- Role Assume events
- Pipeline changes
- Deployment logs

Step 7: Security

- KMS encryption everywhere
- No long-lived credentials
- IAM boundary policies

Step 8: Pipeline Optimizations

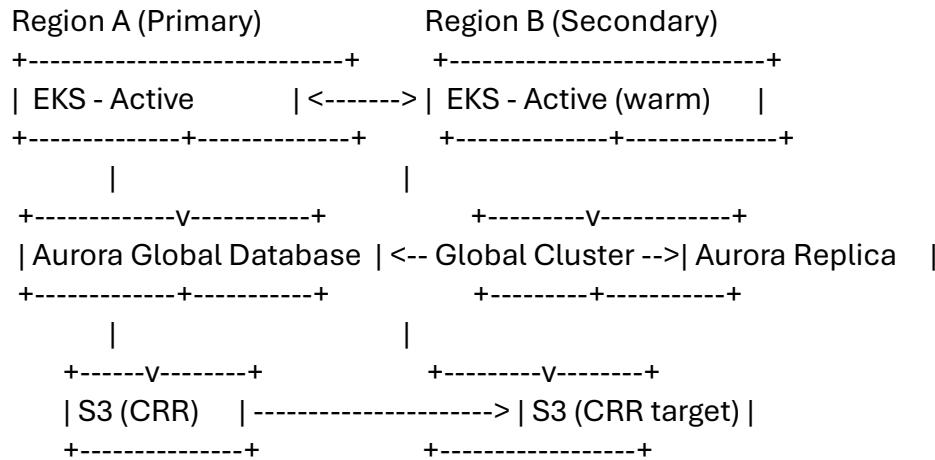
- Parallel stages
- Caching layers
- Canary deployments
- Blue/green for production

ARCHITECTURE 40 — Large-Scale Disaster Recovery Architecture (RPO 0, RTO < 5 Minutes)

Goal:

Achieve near-zero data loss and ultra-fast recovery across two AWS regions.

Diagram



Step-by-Step Architecture

Step 1: Compute Layer — Active/Active or Active/Warm

EKS clusters running in both regions:

- Primary cluster handles traffic
- Secondary cluster scaled to minimum

- Autoscaling kicks in during disaster
-

Step 2: Database Layer — Aurora Global DB

Benefits:

- Sub-second replication
 - Fast failover (<15s)
 - One primary writer, multiple read replicas
-

Step 3: Storage Layer — S3 CRR

Cross-Region Replication:

- Versioning enabled
 - KMS keys replicated
 - Ensures no data loss
-

Step 4: DNS — Global Failover

Route 53 routing policies:

- Health checks
- Failover routing
- Latency-based routing

Failover time:

- 20–40 seconds DNS propagation
 - Combined with Aurora failover = <5 minutes RTO
-

Step 5: Stateless Microservices

Containers store no state → easy failover.

State stored in:

- Aurora
 - Redis Global Datastore
 - S3
-

Step 6: Message Queue Replication

Options:

- Cross-region SQS
 - Kafka cluster replication
 - Kinesis with enhanced fan-out
-

Step 7: Automation

DR scripts perform:

- Failover Aurora
 - Update environment variables
 - Scale EKS nodegroups
 - Shift DNS weights
-

Step 8: Testing

Regular DR drills:

- Chaos engineering
- Simulate region outage
- Validate RTO and RPO