

Batch A — 40 Advanced Scenarios (detailed)

1) EC2 unreachable but system checks pass

Diagnose: Check route/SG/NACL, ENI status, instance system logs.

aws ec2 describe-instance-status, view CloudWatch & EC2 System Log.

Immediate mitigation: Reattach ENI or stop/start instance (non-destructive) if reachable.

Use EC2 Serial Console.

Root fix: Correct SG/NACL/router misconfig or repair corrupted network config.

Prevent: Monitoring on ENI detach events, alerting, runbook.

2) S3 PUT requests spiked 10x

Diagnose: Use CloudTrail to find caller; S3 ServerAccess logs or S3 Storage Lens.

Immediate mitigation: Block the offending IAM user/role; apply bucket policy deny.

Root fix: Fix buggy application, SDK retry logic; remove malicious process.

Prevent: Alerts on unusual PUT/GET rates, IAM permissions review, VPC endpoints.

3) NAT Gateway costs exploded

Diagnose: VPC Flow Logs to find sources/destinations. Check CloudWatch NAT metrics and billing.

Immediate mitigation: Create VPC endpoints for S3/ECR etc. Restrict internet egress.

Root fix: Move frequent S3/ECR calls to endpoints; fix cron jobs/service pulling packages.

Prevent: Daily cost monitors, route auditing, autoscaling of data transfers.

4) Performance degraded after EBS encryption

Diagnose: Check CPU usage (encryption CPU overhead), instance type capabilities (AES instructions).

Immediate mitigation: Move I/O heavy workloads to instances with better crypto performance.

Root fix: Use EBS optimized instances with modern CPUs or provision gp3/io2 with higher IOPS.

Prevent: Benchmark encryption overhead in staging before rollout.

5) RDS latency increase during peak

Diagnose: CloudWatch RDS metrics (CPU, read/write latency, queue depth), slow query log.

Immediate mitigation: Add read replicas / scale vertically; increase connection pool.

Root fix: Query tuning, indexes, offload reads to replicas, optimize schema.

Prevent: Use performance insights, baseline query SLAs, capacity planning.

6) RDS storage skyrockets

Diagnose: Check binary log retention, slow query generation of temp tables, large backups.

SHOW BINARY LOGS; review snapshot schedule.

Immediate mitigation: Purge or archive logs, increase storage if urgent.

Root fix: Adjust log retention, clean temp tables, schedule backups at off-peak times.

Prevent: Alerts at storage thresholds, automated auto-scale.

7) DynamoDB Global Table replication lag

Diagnose: CloudWatch for throttling metrics, check hot keys, CloudWatch Alarms on consume.

Immediate mitigation: Increase WCU/WCU temporarily or switch to on-demand mode.

Root fix: Partition key redesign, implement exponential backoff, use adaptive capacity.

Prevent: Load testing across regions, autoscaling policies.

8) Lambda timeouts

Diagnose: Check CloudWatch logs, X-Ray traces; VPC ENI cold start delays if in VPC.

Immediate mitigation: Increase timeout; provisioned concurrency for critical functions.

Root fix: Reduce cold start by keeping function warm, avoid heavy SDK init on cold start, move out of VPC if possible with VPC endpoints.

Prevent: Instrumentation (X-Ray), SLOs for Lambda durations.

9) CloudFront 504s

Diagnose: CloudFront logs -> origin response timeouts; check origin health.

Immediate mitigation: Increase origin timeout in distribution; scale origin.

Root fix: Address backend performance, tune origin timeouts and caching.

Prevent: Cache dynamic content if possible; use origin shielding.

10) ELB deregistering healthy targets

Diagnose: ELB health check path/port mismatch, check target group health, app process listening port.

Immediate mitigation: Correct health check path or increase timeout; re-register targets.

Root fix: Fix app readiness to respond quickly.

Prevent: Test health checks during deployment; use lifecycle hooks in ASG.

11) Pod running but app dead

Diagnose: kubectl exec into pod, check process list, application logs. Confirm readiness vs liveness probe configuration.

Immediate mitigation: Restart container or pod to restore state.

Root fix: Adjust liveness probe (should kill hung processes) and readiness probe to prevent traffic to unhealthy pods.

Prevent: Use proper probes and graceful shutdowns.

12) kube-proxy high CPU

Diagnose: kubectl get endpoints/services huge counts, iptables rule explosion.

Immediate mitigation: Restart kube-proxy, switch mode to IPVS.

Root fix: Reduce service churn, optimize controllers creating endpoints, use headless services for many clients.

Prevent: Use IPVS mode and monitor service counts.

13) CoreDNS SERVFAIL

Diagnose: kubectl logs -n kube-system coredns ; check upstream DNS, node DNS resolution, CoreDNS config.

Immediate mitigation: Restart CoreDNS pods; scale up replicas.

Root fix: Fix upstream DNS or VPC DNS settings, ensure CoreDNS has enough resources.

Prevent: Alerting on CoreDNS error rates; caching.

14) K8s API server slow

Diagnose: Check control plane metrics, etcd performance, large CRD list or large number of objects.

Immediate mitigation: Scale control plane (managed EKS or control plane autoscaling), reduce watchers.

Root fix: Cleanup unused CRDs, reduce event rate, tune etcd resources.

Prevent: Limit scale of controllers and use leader election.

15) PVC bind delays

Diagnose: StorageClass mismatch, check kubectl describe pvc for events and PVs; AZ mismatch.

Immediate mitigation: Create PV manually if urgent or reschedule pod to correct AZ.

Root fix: Use multi-AZ storage classes or dynamic provisioning with correct parameters.

Prevent: Volume topology awareness and validating YAML.

16) CrashLoopBackOff although code OK

Diagnose: Check container entrypoint, permissions, environment variables, missing files. kubectl logs -p.

Immediate mitigation: Adjust entrypoint or fix config; restart pod.

Root fix: Fix Dockerfile entrypoint, config management, secrets mounting.

Prevent: Integration tests in CI that start containers with actual entrypoints.

17) OOMKilled persists despite more memory

Diagnose: Heap leak, duplicated memory usage, or a process that grows over time. Collect heap/heapdump.

Immediate mitigation: Increase limit short-term; restart pod with HPA.

Root fix: Fix leak in code, tune JVM GC, memory limits/requests properly.

Prevent: Memory profiling, alerts for memory growth trends.

18) Node NotReady due to disk pressure

Diagnose: Node kubectl describe node, check kubelet logs, df -h on node.

Immediate mitigation: Clean large files (logs), cordon and drain node, replace node.

Root fix: Add log rotation, increase disk sizes, move ephemeral storage to EBS.

Prevent: Disk usage alerts, log forwarding to central store.

19) Autoscaler adds nodes but pods still Pending

Diagnose: Check pod events for pod affinity, taints, wrong instance types or insufficient resources per pod (GPU, local SSD).

Immediate mitigation: Add nodegroup with required instance type/labels or remove taints.

Root fix: Ensure cluster-autoscaler supports the nodegroup/labels required, correct resource requests.

Prevent: Pod resource request tuning and nodegroup planning.

20) Downtime during rolling update

Diagnose: Check Deployment strategy maxUnavailable, readiness probes, pod startup time.

Immediate mitigation: Increase replicas, set PDBs, tweak rollingStrategy to smaller batches.

Root fix: Improve startup time, use readiness probes to delay serving.

Prevent: Canary deployments and pre-warming.

21) Terraform state lost

Diagnose: Check S3 versioning, backend misconfiguration; list state file versions.

Immediate mitigation: Restore state from S3 versioning or backup, use terraform import to reconstruct resources if needed.

Root fix: Strict backend management, enable state locking and versioning.

Prevent: Enforce state snapshots, automated backups.

22) Two devs applied Terraform simultaneously

Diagnose: Locking failure or overwritten state. Check DynamoDB lock table if used.

Immediate mitigation: Reconcile resource changes manually, run terraform refresh, import missing resources.

Root fix: Enforce remote state locking (DynamoDB) and PR process for infra changes.

Prevent: Pre-apply approvals, CI gated infrastructure changes.

23) CI pipeline runs but deploy not reaching cluster

Diagnose: Check kubeconfig used by pipeline, RBAC, network connectivity from runner.

Immediate mitigation: Verify runner credentials and context; redeploy with correct kubeconfig.

Root fix: Secure and centrally manage service accounts; rotate tokens.

Prevent: Pipeline preflight checks confirm cluster connectivity.

24) Helm upgrade failed, cluster unstable

Diagnose: Partial resources applied, failed CRDs, hooks stuck. helm history and helm status to inspect.

Immediate mitigation: helm rollback to last stable release. Clean orphan resources.

Root fix: Validate manifests, add precheck hooks and dry runs.

Prevent: Use helm test and CI linting, add resource quotas.

25) Applied Terraform to production accidentally

Diagnose: Pause pipelines, review change plan output.

Immediate mitigation: Stop the apply if possible; roll back to last known state; restore from backups.

Root fix: Separate state files per environment, require manual approvals for prod.

Prevent: Protect production workspace with approval gates and limited exec permissions.

26) Docker build 10 minutes

Diagnose: Large base images, many layers, dependencies reinstalled each build.

Immediate mitigation: Use caching, build on bigger hardware for now.

Root fix: Multi-stage builds, smaller base images (alpine/distroless), cache dependencies.

Prevent: Use CI cache and prebuilt base images.

27) Jenkins agent containers restart often

Diagnose: Check Jenkins logs & agent container logs for OOM or disk full.

Immediate mitigation: Increase agent resources/ephemeral storage.

Root fix: Clean workspace, use ephemeral agents with enough disk size; limit concurrent builds.

Prevent: Monitor disk, set quotas, ephemeral agent lifecycle.

28) Secrets pushed to Git

Diagnose: Identify commit(s) using git log and git grep.

Immediate mitigation: Rotate compromised secrets immediately. Remove secrets from repo using git filter-repo or BFG, push forced history rewrite.

Root fix: Move to Secrets Manager/SSM Parameter Store and enforce pre-commit hooks.

Prevent: Secret scanning in CI, deny commits with secrets, rotate secrets regularly.

29) Works in staging but fails prod

Diagnose: Compare env vars, IAM roles, VPC/network differences, feature flags.

Immediate mitigation: Recreate prod with same configs as staging for controlled tests; revert to previous stable release.

Root fix: Environment parity, config as code, use the same secrets and permission models.

Prevent: Promote artifacts, not builds; blue/green strategy.

30) Unexpected auto-reverts

Diagnose: Auto rollback policies triggered by health checks or custom monitors. Check pipeline and deployment health events.

Immediate mitigation: Identify failing checks and either fix the reason or temporarily disable auto-revert to inspect.

Root fix: Fix underlying instability causing health checks to fail (performance, smoke tests).

Prevent: Tighter pre-prod validation, canary waits.

31) 5xx errors but CPU/memory normal

Diagnose: Check DB connection pools, downstream services, thread dumps, timeouts.

Use tracing (Jaeger/X-Ray).

Immediate mitigation: Restart flaky downstream services, increase DB pool or connections temporarily.

Root fix: Fix connection leaks, tune timeouts, increase retry/backoff.

Prevent: Instrument and set SLOs for downstream calls.

32) Daily latency spikes at same time

Diagnose: Investigate cron jobs/backup snapshots, DB compactions, garbage collection. Correlate with cloudwatch metrics.

Immediate mitigation: Reschedule heavy jobs to off-peak.

Root fix: Stagger jobs, optimize backups.

Prevent: Run capacity planning and schedule maintenance windows.

33) Memory leak in production

Diagnose: Capture heap dump, analyze with profiler (YourKit/VisualVM), trace leak path.

Immediate mitigation: Restart pods in a rolling fashion to clear memory.

Root fix: Fix code causing leak, add stream/iterator closures.

Prevent: Add memory growth alarms, periodic Canary restarts.

34) DB connection exhaustion across services

Diagnose: Check max_connections, connection pool usage per service, long transactions.

Immediate mitigation: Increase DB connection limit or add connection pooling (pgbouncer).

Root fix: Implement connection pooling, reduce per-request connections, tune pool sizes.

Prevent: Monitor connection counts, circuit breaker for DB.

35) EBS IOPS exhaustion

Diagnose: CloudWatch EBS metrics, identify high I/O disks, check instance type for EBS optimization.

Immediate mitigation: Move hotspot I/O to provisioned IO (io2), use caching (ElastiCache), or add sharding.

Root fix: Restructure workload, use higher IOPS volumes or scale horizontally.

Prevent: Configure appropriate IOPS and monitor.

36) Error budget exceeded

Diagnose: Review SLO metrics, determine common failure modes.

Immediate mitigation: Freeze releases to production; prioritize reliability work.

Root fix: Fix the top contributors to errors.

Prevent: Invest part of sprint time for reliability and safety for releases.

37) Client sees slowness but server shows healthy

Diagnose: Test from client location (traceroute), CDN edge health, Route53, DNS TTLs, network latency.

Immediate mitigation: Use region-specific endpoints or improve CDN caching.

Root fix: Increase edge caching, check peering issues with ISP.

Prevent: Synthetic monitoring from multiple geographies.

38) Kubernetes network jitter rising

Diagnose: Check CNI plugin metrics, node saturation, iptables complexity.

Immediate mitigation: Move high throughput pods to separate nodes, restart CNI.

Root fix: Switch to high-performance CNI/IPVS, upgrade kernel and CNI plugin.

Prevent: Network QoS and node sizing, isolate noisy neighbors.

39) Logs increased 20x

Diagnose: Identify source service enabling debug, look for log spam or error loops.

Immediate mitigation: Temporarily reduce log level for that service, clean up disk space.

Root fix: Fix the underlying error causing log spam.

Prevent: Set log rate limits and structured logging with sampling.

40) Incident repeats despite fixes

Diagnose: Review RCA, check whether root cause was fixed or only symptoms. Examine change control & test coverage.

Immediate mitigation: Implement temporary controls (circuit breakers, throttling) to stop recurrence.

Root fix: Comprehensive root cause elimination (code fix + infra change + process).

Prevent: Add regression tests, automation to prevent recurrence, documented runbooks.

Batch B — 40 More Advanced Scenarios (detailed answers)

(These correspond to the second set of 40 scenarios I gave earlier — answers follow the same Diagnose → Immediate → Root-Cause → Prevent pattern.)

41) VPC IP exhaustion because of EKS nodes

Diagnose: aws ec2 describe-subnets to check allocated IPs per subnet. Count pod IP allocation using AWS VPC CNI.

Immediate mitigation: Add new subnets with larger CIDR and move nodegroups, or enable WARM_IP_TARGET to reduce pod-per-node IP.

Root fix: Re-architect with larger CIDRs or use secondary CIDR ranges/VPC peering or use CNI overlay (kube-ovn/calico).

Prevent: Plan IP addressing for scale, use pod IP reuse features.

42) ALB high target response time but EC2 CPU normal

Diagnose: Check network latency, backend thread pools, unhealthy dependency (DB), disk I/O, and connection reuse. Use X-Ray/tracing.

Immediate mitigation: Scale backend, add cache, increase connection pool.

Root fix: Identify blocked threads or slow downstream service; optimize queries or code.

Prevent: Add distributed tracing and alert on tail latency.

43) RDS failover too slow

Diagnose: RDS failover time depends on engine and replication lag; check CloudWatch failover metrics.

Immediate mitigation: Use Multi-AZ provisioned with synchronous replication (Aurora is faster).

Root fix: Migrate to Aurora Global DB or use read-replicas with promoted failover and fast DNS TTLs.

Prevent: Test failovers regularly and tune maintenance windows.

44) S3 replication causing duplicate data

Diagnose: Check replication configuration, multiple replication rules, or retry loops from source.

Immediate mitigation: Stop the process causing duplicate writes, delete duplicates if safe.

Root fix: Fix replication rules, ensure object versioning and idempotent writes.

Prevent: S3 events/metrics and deduplication keys.

45) Lambda cold starts

Diagnose: Check duration histograms, factor in VPC ENI attach time, package size, runtime.

Immediate mitigation: Use provisioned concurrency for critical functions.

Root fix: Reduce package size, lazy-init modules, avoid VPC when not needed, use VPC endpoints.

Prevent: Warmers, metrics on cold start frequency.

46) NAT Gateway bill — repeated (detailed)

(See #3) additionally inspect ECR pulls and periodic updates from many instances. Use CloudWatch Contributor Insights to find top talkers.

47) ENI attach failures during ASG scaling

Diagnose: Check EC2 instance limit for ENIs per instance type, IAM permissions for CNI to attach ENIs.

Immediate mitigation: Use instance types with higher ENI quotas or change pod density.

Root fix: Update instance types or CNI plugin settings, ensure amazon-vpc-cni has permissions.

Prevent: Capacity planning for ENIs and pod density.

48) Intermittent packet drops across VPCs

Diagnose: VPC peering / Transit Gateway route table misconfigurations, SG NACL intermittent rules, network hardware events.

Immediate mitigation: Reconfigure route tables or failover to alternative path, open necessary ports.

Root fix: Fix route or TGW attachments, increase MTU alignment across path.

Prevent: Redundant network paths and monitoring.

49) Migrate VPC CIDR without downtime

Diagnose: This is a delicate migration. Plan for new VPC, transit gateway peering, route propagation.

Immediate mitigation: Create new VPC with desired CIDR, set up VPC peering or TGW and gradually move resources via new subnets (use elastic IPs, DNS cutover).

Root fix: Full migration plan.

Prevent: Plan VPC allocation ahead, use secondary CIDRs for expansion.

50) CloudFront caching not working for dynamic content

Diagnose: Check Cache-Control headers, query string forwarding, cookies and CacheBehavior.

Immediate mitigation: Adjust Cache-Control to allow caching; configure behaviors for specific paths.

Root fix: Update application to set proper caching headers and TTLs for dynamic/static separation.

Prevent: Test caching behavior in staging.

51) Pods stuck Terminating

Diagnose: Finalizers on resource, webhook blocking termination, PV unmount issues. kubectl get pod -o yaml to inspect finalizers.

Immediate mitigation: Remove finalizer if safe; force delete as last resort.

Root fix: Fix webhook handlers or ensure PV unmounts properly.

Prevent: Build graceful shutdown handlers and test.

52) CoreDNS pods restarting — more detail

Diagnose: Check CPU/memory, upstream DNS unreachable, recursive loops.

Immediate mitigation: Scale CoreDNS replicas; point CoreDNS to stable resolvers.

Root fix: Fix networking issues or upstream DNS failure.

Prevent: Resource limits and readiness probes for CoreDNS.

53) HPA not triggering

Diagnose: Check metrics server is running (metrics-server), proper metrics configured; check target utilization and resource requests.

Immediate mitigation: Manually scale if necessary.

Root fix: Ensure metrics server is healthy or use external metrics adapter.

Prevent: Validate metric availability in pre-prod.

54) Cluster Autoscaler adds nodes but pods Pending — again

(See #19) additional angle: pods may have nodeSelector requiring GPU or taints. Add matching nodegroup.

55) Worker nodes failing to join cluster

Diagnose: Bootstrap script errors, kubelet certificate issues, wrong AMI, IAM role or security group blocking.

Immediate mitigation: Inspect cloud-init logs on node, replace node.
Root fix: Fix AMI or bootstrap process.
Prevent: Bake AMIs with required kubelet components.

56) App slow after rolling update

Diagnose: Check new version resource usage, compatibility changes, database migration steps causing locks.
Immediate mitigation: Rollback to previous version.
Root fix: Improve migration strategy and perform canary testing.
Prevent: Blue/green and feature flags.

57) Deployment stuck at 0/5 updated pods

Diagnose: ImagePullBackOff, insufficient resources, failed init containers. kubectl describe rs and pods.
Immediate mitigation: Fix image path/credentials or add resources.
Root fix: CI ensures images are pushed to correct registry with tags.
Prevent: Pre-deploy smoke test.

58) OOMKilled despite increase — deeper

Diagnose: Host-level memory fragmentation or swap issues, cgroup limits, memory leak persists.
Immediate mitigation: Add swap or add nodes, restart pods.
Root fix: Profiling the app; fix memory leaks.
Prevent: Memory pressure monitoring.

59) Noisy neighbor in shared cluster

Diagnose: Identify pod with high resource consumption using kubectl top pod and node metrics.
Immediate mitigation: Apply quotas and limits; cordon node and move noisy pods.
Root fix: Enforce resource requests & limits, use namespaces with ResourceQuotas.
Prevent: Pod QoS classes and burstable handling.

60) CSI driver PVC mount failures

Diagnose: Check CSI driver logs, CSI controller and node plugin status, driver compatibility with K8s version.
Immediate mitigation: Restart CSI pods, re-provision PV.
Root fix: Update CSI driver, fix IAM roles for driver.
Prevent: Compatibility testing and readiness probes for CSI.

61) Terraform corrupted state (more detail)

Diagnose: Identify state version in S3, check lock.
Immediate mitigation: Restore previous version from S3, run terraform state rm to remove problematic resources.

Root fix: Add strict locking and pre-apply reviews.

Prevent: Enforce CI validation and backups.

62) Terraform plan shows large unwanted changes

Diagnose: Compare local vs remote state, provider upgrades, resource rename.

Immediate mitigation: Pin provider version; import resources back into state if they were created outside TF.

Root fix: Use terraform import carefully and ensure state refresh.

Prevent: Terraform runbook and provider pinning.

63) Pipeline runs but deployment not reaching cluster (revisit)

(Same as #23) include checking Kubernetes RBAC and ServiceAccount token TTL.

64) Helm upgrade failed — cluster unstable (revisit)

(See #24) add helm hooks idempotency.

65) Applied TF in prod mistakenly — further actions

(See #25) and implement immediate freeze and postmortem.

66) Optimize Docker build time (more)

Diagnose: Repeated apt-get/ package installs, no build cache.

Immediate mitigation: Use build cache retention; use buildkit.

Root fix: Separate dependencies into separate layers and use cache from registry.

Prevent: CI caching and reproducible builds.

67) Jenkins agent container restarts — additional

Include ephemeral storage pressure from logs; use docker system prune cron.

68) Secrets in Git — legal/compliance actions

Also audit access logs and notify security team; rotate secrets across services and dependent CI.

69) Staging vs Prod mismatch — more

Add feature toggles and ensure IaC templates identical; use environment matrix.

70) Auto reverts — more depth

Investigate health checks flapping thresholds, circuit breakers; perhaps overzealous monitoring triggers.

71) 5xx errors — deep trace

Use distributed tracing to locate service causing 5xx; inspect rate limiter behavior and queue lengths.

72) Nightly latency spike — further

Check snapshot/backup windows for RDS/Aurora; also check S3 lifecycle jobs and repackaging jobs.

73) Memory leak in prod — extra

If language is Java, collect GC logs and perform heap dump analysis; if Python, check reference cycles.

74) DB connection spikes across microservices — additional

Enforce connection pooling libraries (HikariCP), tune max_connections per app and DB.

75) Disk IOPS exhaustion — more

Consider read replicas for read-heavy workloads, upgrade to provisioned IOPS, or use local NVMe for caches.

76) Error budget exceeded — extra

Communicate to business stakeholders and schedule reliability sprint.

77) Client slowness with server healthy — more

Check CDN edge misconfigurations, large TLS handshakes, renegotiation storms.

78) K8s network jitter — more

Examine CNI plugin version, tune MTU, and check node-level CPU steal affecting network stacks.

79) Logs increase 10x — additional steps

Set up adaptive logging where high severity levels are retained but debug-level logs sampled.

80) Repeat incident — enforce closure

Define success criteria for RCA closure and require proof-of-fix (automated test) before closing.